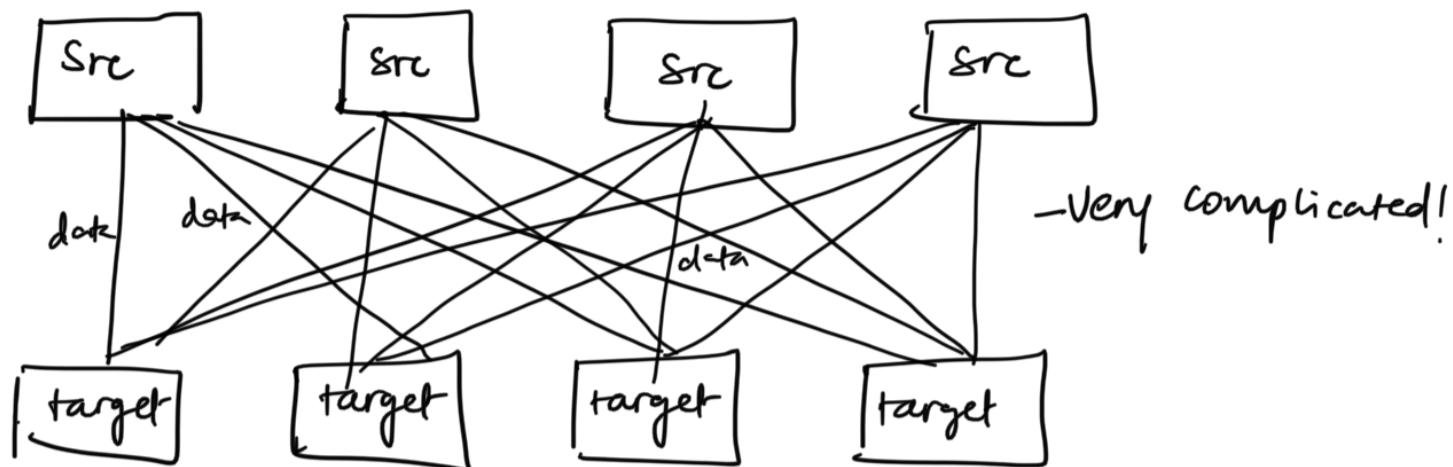


## Apache Kafka

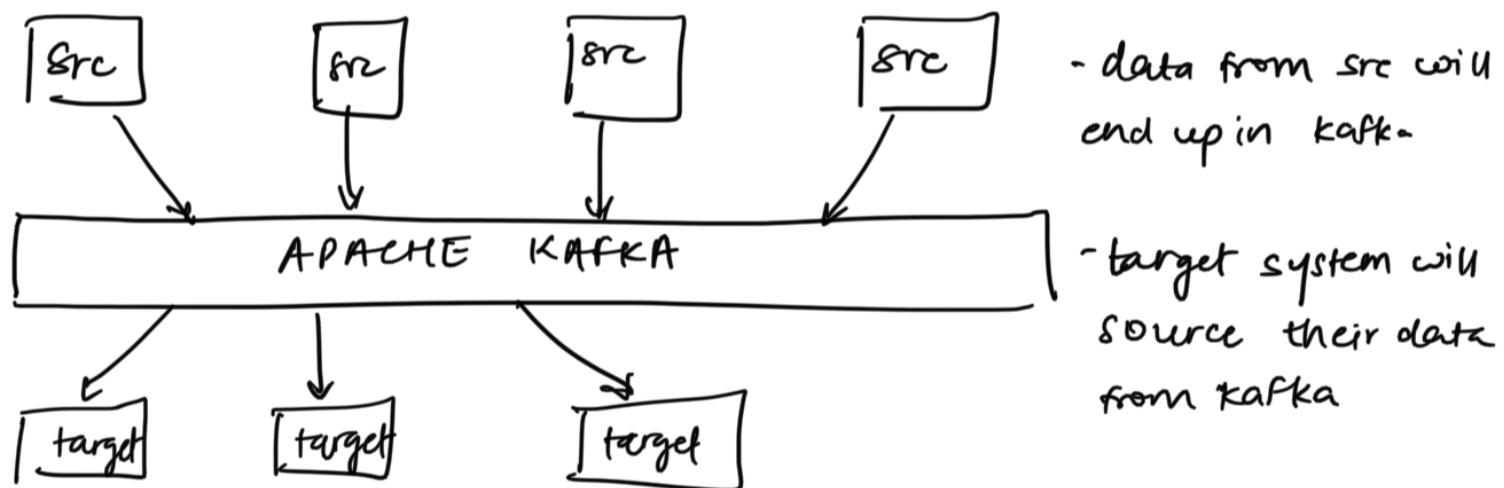
- Previously -



- If you 4 src and 6 target you need  $6 \times 4 = 24$  integrations
- Each integration has a lot of difficulties
  - Protocol, Data format, Data schema & evolution.
- Each source system will have an increased load from the connections.

### \* Apache Kafka -

- allows to decouple data stream and the system.
- software platform based on distributed streaming process. It is a publish-subscribe messaging system which lets exchanging of data between applications, servers and processes as well.



### \* Why Apache Kafka??

- Created by LinkedIn, now Open Source Project mainly maintained by Confluent
- Distributed, resilient architecture, fault tolerant
- Horizontal Scalability
  - can scale to 100's of brokers
  - can scale to millions of messages per second
- High performance (latency less than 10ms) - real time
- Used by 2000+ firms, 35% of fortune 500 - Netflix, Walmart, Uber etc.

### \* Use cases:

- Messaging System
- Activity tracking
- Gather metrics from many different locations

- Application logs gathering
- Stream processing (kafka streams API or spark)
- Decoupling of system dependencies.
- Integration with Spark, flink, storm, Hadoop and many other Big Data Technologies.

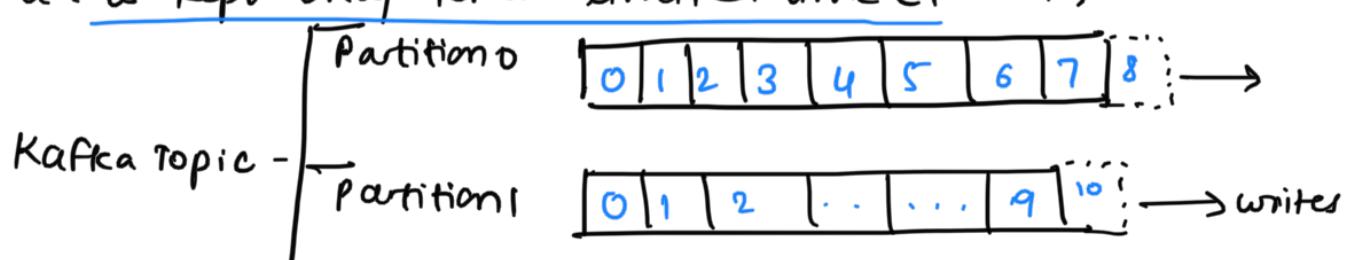
### Apache Kafka's - core API's

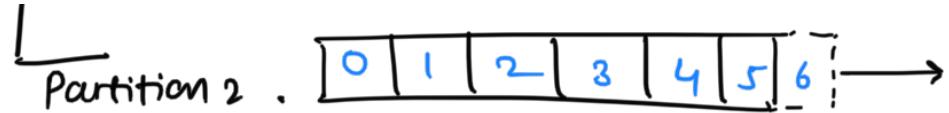
- 1) Producer API - allows application to publish streams of records to one or more topics.
- 2) Consumer API - allows an application to subscribe one or more topics and process the stream of records produced to them.
- 3) Streams API - allows an application to effectively transform input streams to output streams.
  - permits an application to act as stream processor which consumes an input stream from one or more topics and produce an output stream to one or more output topics.
- 4) Connector API - This API executes the reusable producer and consumer API with existing data systems or applications.

### Topics, Partition, Offset.

Topics: a particular stream of data

- Similar to table in a database (without all constraints)
- you can have as many topics as you want.
- A topic is identified by its name
- Topics are split into partitions
  - Each partition is ordered
  - Each message within a partition gets an incremental id, called offset.
  - offset only have a meaning for a specific partition  
eg: offset 3 in partition 0 doesn't represent the same data as offset 3 in partition 1.
  - Order is guaranteed only within a partition (not across partitions)
  - Data is kept only for a limited time (1 week)





- Once data is written to a partition it can't be changed.
- Data is assigned randomly to a partition unless a key is provided

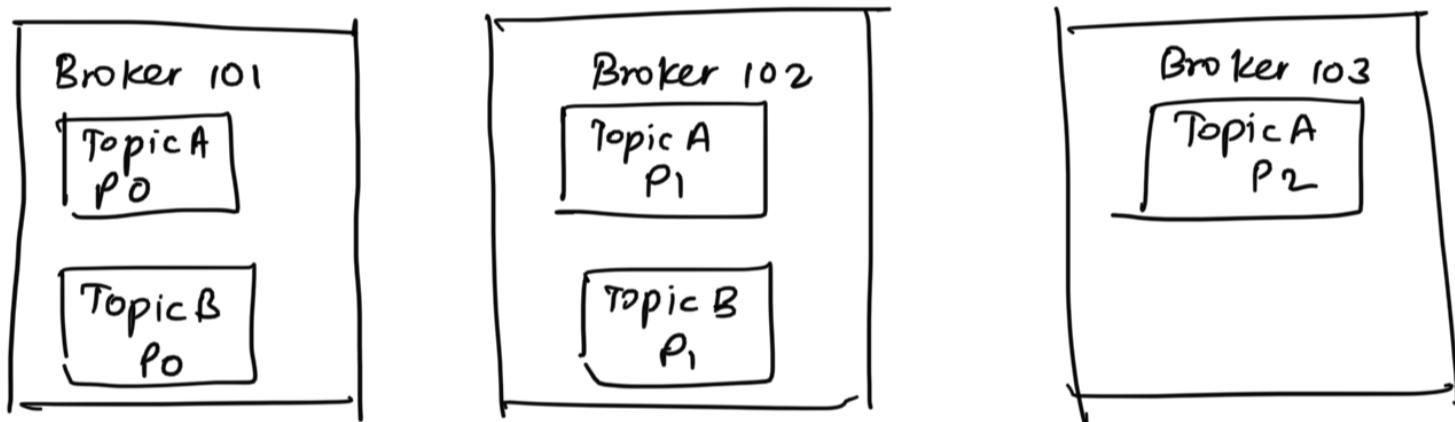
### Brokers

- A Kafka cluster is comprised of multiple brokers (servers).
- Each broker is identified with its ID (integer)
- Each broker contains certain topic partitions
- \* After connecting to any broker (called a bootstrap broker) you will be connected to entire cluster.

### Brokers and Topics

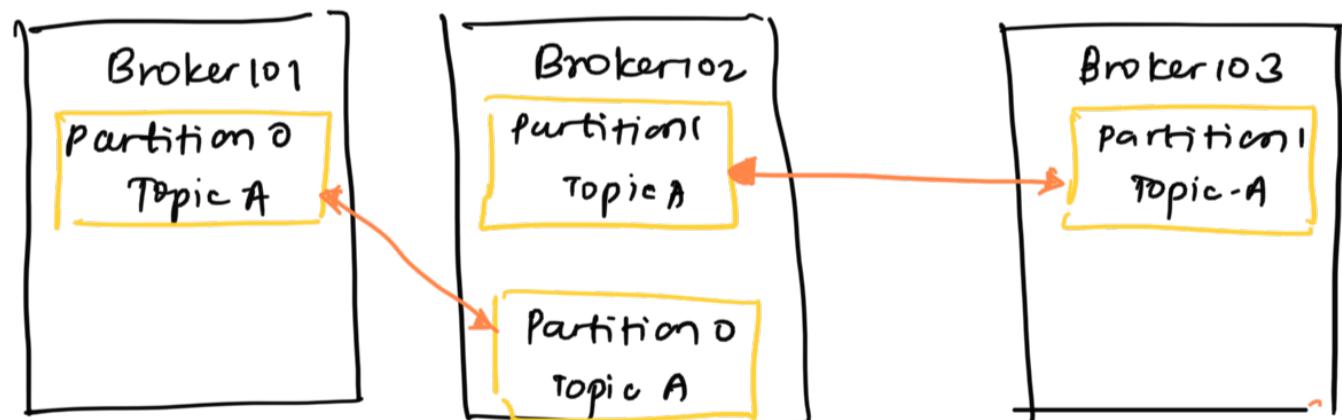
Topic A - 3 partitions (P<sub>0</sub>, P<sub>1</sub>, P<sub>2</sub>)

Topic B - 2 partitions (P<sub>0</sub>, P<sub>1</sub>)



### Topic Replication Factor

- >1 (usually b/w 2 and 3)
- This way if one broker is down, another broker can serve the data.  
eg. Topic A - 2 partitions, replication factor 2



### Concept of leader for a partition

- At any time only ONE broker can be a leader for a given partition.
- Only that leader can receive and serve data for a partition.
- The other brokers will synchronize the data
- Therefore each partition has one leader and multiple in sync replicas.  
(ISR)

## Producer.

- writes data to topics (which is made of partitions)
- producer automatically know to which broker and partition to write to.
- In case of broker failures, producers will automatically recover
- producer will load balance in absence of keys. It will divide data between different partitions.

Producer can receive acknowledgement of data writes.

acks=0: producer won't wait for ack (possible data loss).

acks=1: will wait for leader ack. (limited data loss).

acks=all: leader + replicas ack (no data loss)

## Message keys:

- producer can choose to send a key with the message (string, number...)
- If key=null, data is sent round robin. (101→102→103...)
- If key is sent, then all messages for that key will go to the same partition always.
- A key is basically sent if you need message ordering for a specific field. (eg. truck-id)

## Consumer.

- read data from a topic (identified by name)
- consumers know which broker to read from.
- In case of broker failures, consumers know how to recover.
- Data is read in order within each partitions.

## consumer groups:

- consumer reads data in consumer groups.
- each consumer within a group reads from exclusive partitions.
- If you have more consumers than partitions, some will be inactive.
- Consumers know where to read from. they will automatically use a group co-ordinator and a consumer co-ordinator to assign consumers to a partition.

## Consumer offsets:

- Kafka stores offsets at which a consumer group has been reading.
- The offsets committed live in a kafka topic named -- consumer-offsets.
- When a consumer in a group has processed data received from Kafka it should be committing the offsets.
- If a consumer dies it will be able to read back from where it left off, thanks to committed consumer offsets.

Consumers choose when to commit offsets.

There are 3 delivery semantics -

At most once - offsets are committed as soon as the message is received.

- If processing goes wrong, the message will be lost (it won't be read again)

At least once - offsets are committed after the message is processed

- if processing goes wrong, the message will be read again.

- This can result in duplicate processing of msgs.  
So make sure your processing is idempotent  
(ie - processing again the messages won't impact your systems).

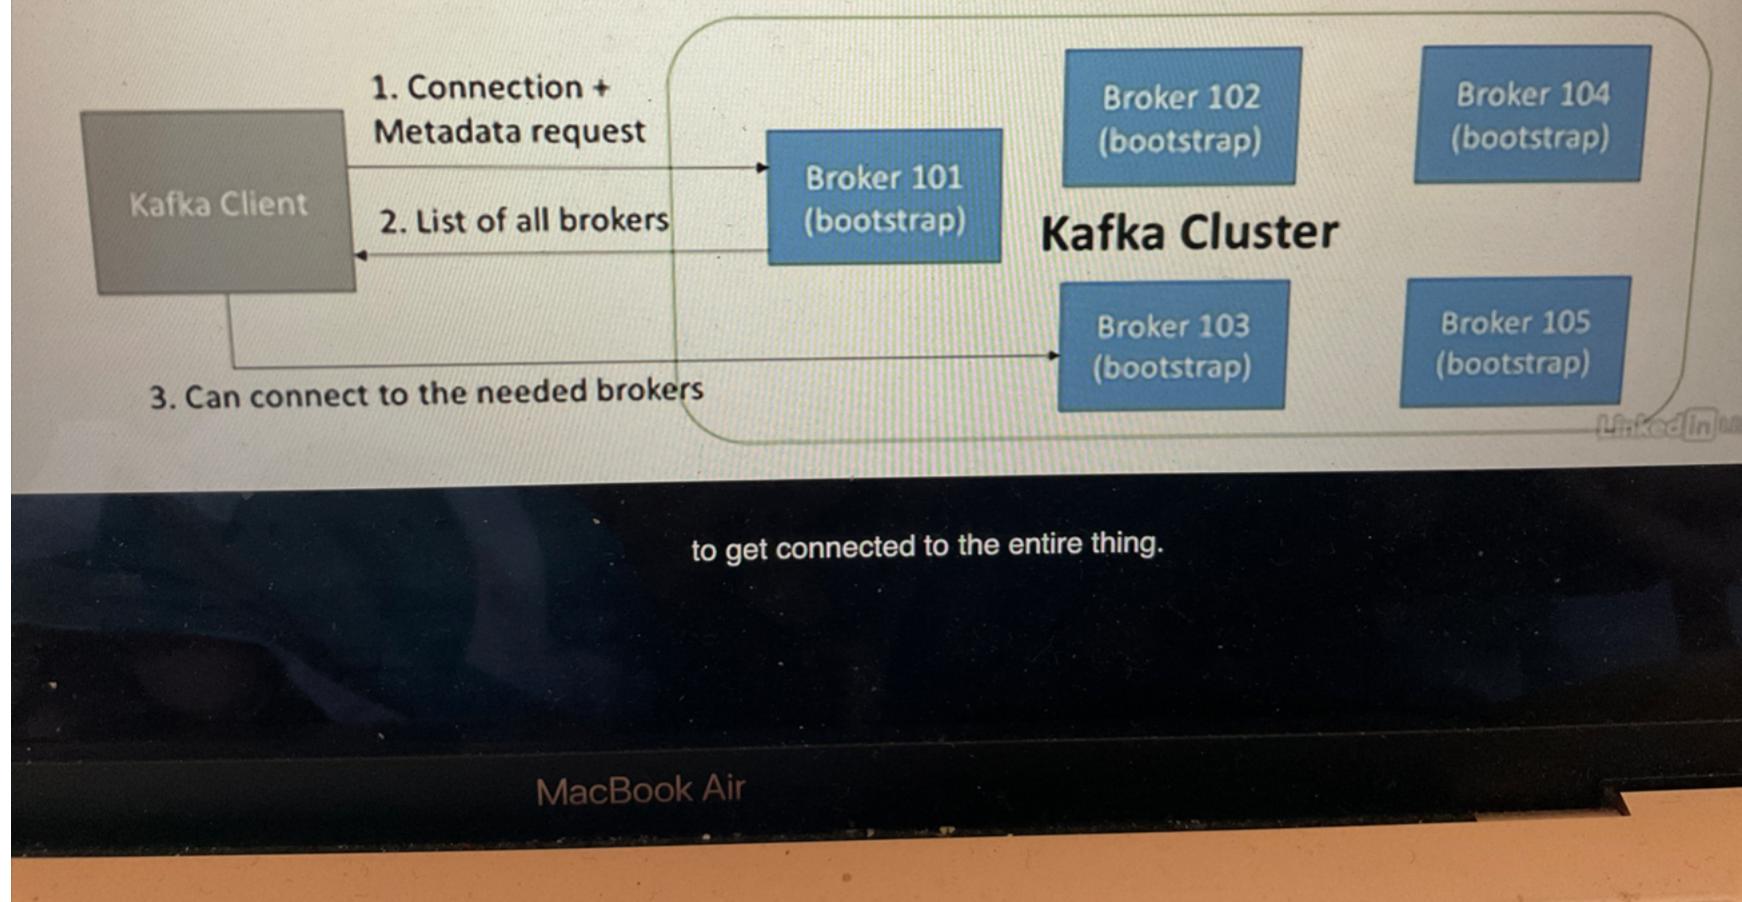
Exactly once: can be achieved for kafka  $\Rightarrow$  kafka workflows using kafka streams API.

For kafka  $\Rightarrow$  external system workflows, use an idempotent consumer.

### Kafka Broker Discovery

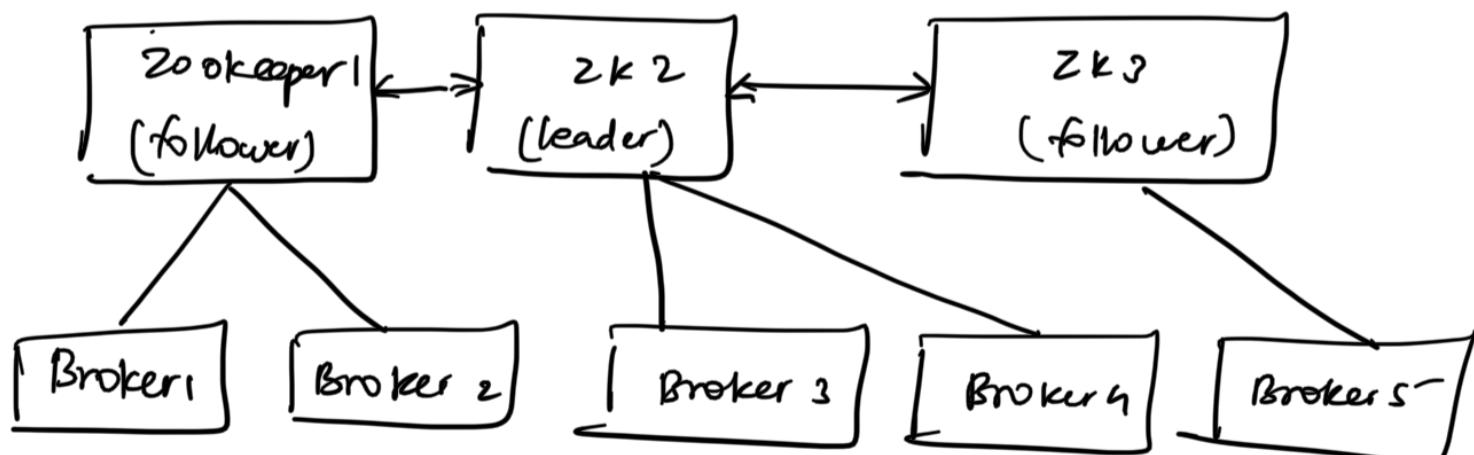
- Each broker is also called a bootstrap broker
- This means you only need to connect to one broker and you will be connected to entire cluster.
- Each broker knows about all brokers, topics and partitions (metadata).

- Every Kafka broker is also called a “bootstrap server”
- That means that **you only need to connect to one broker**, and you will be connected to the entire cluster.
- Each broker knows about all brokers, topics and partitions (metadata)



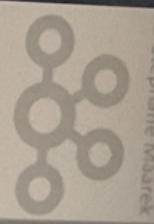
### Zookeeper:

- manages brokers (keeps a list of them)
- helps in performing leader elections for partitions.
- Sends notifications to kafka in case of changes (e.g. new topic, broker dies, broker comes up etc)
- Kafka can't work without Zookeeper.
- Zookeeper by design operates on an odd number of servers (3, 5, 7).
- Zookeeper has a leader (handles writes) the rest of the servers are followers.
- Zookeeper does not store consumer offset with kafka (v0.10)



### Kafka Guarantees:

# Kafka Guarantees



- Messages are appended to a topic-partition in the order they are sent
- Consumers read messages in the order stored in a topic-partition
- With a replication factor of N, producers and consumers can tolerate up to N-1 brokers being down
- This is why a replication factor of 3 is a good idea:
  - Allows for one broker to be taken down for maintenance
  - Allows for another broker to be taken down unexpectedly
- As long as the number of partitions remains constant for a topic (no new partitions), the same key will always go to the same partition

LinkedIn 418 active

you'll remember them, I promise,

## Theory Round up:

418 active

# Theory Roundup

## We've looked at all the Kafka concepts



Stephanie Marek

