

Consumer Poll Behavior

* Kafka consumers have a poll model that means they will ask data from kafka.

push model - server pushes data & consumer waits.

poll model - allows consumers to control where in the log they want to consume, how fast and gives them the ability to replay events.

* fetch.min.bytes (default: 1):

- controls how much data you want to pull at least on each request.
- helps improving throughput and decreasing request number.
- At the cost of latency.

* Max.poll.records (default: 500):

- controls how many records to receive per poll request
- Increase if your messages are very small and have a lot of available RAM.
- Good to monitor how many records are polled per request.

Max.partitions.fetch.bytes (default 1MB):

- max data returned by the broker per partition.
- If you read from 100 partitions, you'll need a lot of RAM

fetch.max.bytes (default 50MB)

- max data returned for each fetch request (covers multiple partitions).
- the consumer performs multiple fetches in parallel.

change these settings only if your consumer maxes out on throughput already