

- SQL -

Wildcard (*)

- used to select all available columns from a table.
- This effect applies to all tables the query accesses through JOIN clauses.
- To select all values from a specific table, the wildcard character can be applied to table with dot notation

SELECT

```
Employee.* , Department.Name  
FROM  
    Employee JOIN Department  
ON  
    Departments.ID = Employee.DeptID
```

Warnings against use -

- Use of * is avoided in production code cause it can cause following problems.
 - * 1. Excess IO, network load, memory use and so on
 - * 2. further excess IO if DB needs to spool internal results to disk as a part of processing query
 - * 3. Potential ambiguity in column names if columns are added to tables and views later on.

when to use -

- to perform manual queries against db for investigation and prototype work.

Alias -

```
{ Select fname as "firstName"      → all versions of SQL.  
    select fname as "firstname"  
          mname as 'middle'  
          lname as [last name] }   { -MySQL ("", . . . , [])  
    from Employees
```

If alias has a single word that is not a reserved word we can write it without single quote, double quote or []

eg

```
Select fname as firstName from employees;
```

```
Select fullname = firstName + ' ' + lastName from emp → MySQL
```

```
Select firstName + ' ' + lastName AS fullName →
```

'As' - standard rather than '='

If you need to use reserved words as alias you can use brackets or quote to escape

Select fname as "SELECT"
MNAME as "FROM" → all versions
from emp

Select fname as "SELECT"
lname as "FROM" → MySQL.
LNAME as WHERE

Also, column names may be used any of final clauses of same query

e.g. SELECT
 fname as firstname,
 FROM
 emp
 ORDER BY
 firstname DESC.

Select individual columns

Select name, id, contact from emp

* Columns are returned in order in which they appear in SELECT clause

Select specific number of records

SQL 2008 Standard -

Select id, name from Product

order by price desc

OFFSET 5 ROWS → skip 5 rows

FETCH FIRST 10 ROWS ONLY → limit to 10 rows in off

SQL Server & MS Access

Select TOP 10 id, name from Products → TOP operates after where clause

MySQL, Postgres

Select id, name from Product
limit 10.

Oracle (ROWNUM)

Select id, name from xyz

where ROWNUM <= 10
Order by price desc

→ ROWNUM works as a part of where clause so if other condns do not exist in specified no. of rows you will get 0 results.

Select with condition

select name where id > 10

Select with CASE

When results need some logic applied on the fly one can use CASE statement to implement it.

```
SELECT CASE WHEN col1 < 50 THEN 'under' ELSE 'over'  
END threshold.
```

Also can be chained

SELECT

```
CASE WHEN col1 < 50 THEN 'under'  
WHEN col1 > 50 AND col1 < 100 THEN 'between'  
ELSE "over"  
END threshold
```

FROM tablename

You can also have case inside case.

Select columns which are named after reserved keywords

SELECT

"ORDER",

ID

FROM orders.

Note - This makes column name case sensitive.

SELECT [order], ID from orders. → SQL Server

SELECT 'Order' id from orders → MySQL, MariaDB

Selecting with table alias.

SELECT e.fname, e.lname FROM Employees e

* Note once you use alias, you cannot use canonical table name anymore i.e

Select e.fname, Employee.id from employee e
will throw an error.

Selecting with more than 1 condition

select name from persons where gender = 'M' AND age > 20;

||

OR age > 20;

can be combined

SELECT name from persons WHERE (gender = 'M' AND age < 20)

OR

(gender = 'F' AND age > 20);

Selecting without locking the table

sometimes when tables are used mostly (or only) for reads, indexing does not help anymore and every little bit counts, one might use selects without LOCK to improve performance

SQL Server

SELECT * FROM Tablename WITH (nolock)

MySQL.

SET SESSION TRANSACTION ISOLATION LEVEL READ UNCOMMITTED;

Select * from tablename

SET SESSION TRANSACTION ISOLATION LEVEL REPEATABLE READ;

Oracle

SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED;

SELECT * FROM Tablename;

DB2

SELECT * FROM TABLNAME WITH UR.

Uncommitted read.

Selecting with Aggregate functions

SELECT AVG(Salary) FROM Employees WHERE DeptId = 1;

can be combined with where and group by clause too

MIN()

MAX()

AVG()

Count(), Count(*) (null values are not counted)

SUM(Salary)

Select with condition of multiple values from column.

SELECT * FROM Cars WHERE status IN ('Wait', 'Work')
" "

Select * from Cars where status = 'Wait' OR status = 'Work'.

NOT IN ()

DISTINCT keyword.

Select Count(DISTINCT DeptId) FROM Emp

Get Aggregated Results for Row Group.

Counting rows based on a specific column value

SELECT category, COUNT(*) AS item-count
FROM item
GROUP BY category

Getting average income by department

SELECT dept, Avg(income)
FROM employees
GROUP BY dept

* The important thing is to select only columns specified in the GROUP BY clause or used with aggregate functions.

WHERE clause can be used with GROUP BY, but WHERE filters out records before any grouping is done

Select dept, Avg(income)
FROM employees
Where dept <= 'ACCOUNTING'
Group by dept;

If you need to filter result after grouping use HAVING.

Select dept, avg(income)
FROM employees
Group by dept
Having avg(income) > 1000;

selection with sorted results.

Select * from emp order by lname , fname DESC.
↳ first sort by lname (ASC) then by fname (DESC)
default - ASC.

- Select id, name, number from Employees order by 3;
 ^
 Column number

Select id, fname, lname, number from Employees | fname

SELECT name, phone, number FROM employees
ORDER BY CASE WHEN lname = "Jones" THEN 0 ELSE 1 ASC

↓

This will sort to have all records with lname Jones at top

Selecting with null

Select Name FROM customers where number IS NULL
→ IS NOT NULL

Select distinct (unique values only).

Select DISTINCT continentcode FROM countries;

Select from multiple tables.

Select * FROM table1, table2 → cross product

Select table1.col1, table2.col2 FROM table1, table2

GROUP BY

Result of a select query can be grouped by one or more columns using GROUP BY statement: all results with the same value in the grouped columns are aggregated together.

- Group by can be used in conjunction with aggregation functions using HAVING statement to define how non grouped columns are aggregated

Select emplid, sum(monthlySalary)
FROM Employee
GROUP BY emplid

filter Group By using HAVING clause.

Select a.id, b.name, count(*) BooksWritten
From Bookauthors ba
INNER JOIN
Authors a
ON a.id = ba.authorid
GROUP BY
aid, a.Name
HAVING count(*) > 1 → having BooksWritten > 1
→ return all authors who
create more than one book.

Use Group by to count number of rows for each unique entry in given column.

Eg Name House

Arya Stark

Sansa	stark
Cersei	Lannister
Mara	Greyjoy

```

select house , count(*) number_of_westerosians
from westerosians
group by house
order by number_of_westerosians DESC
    
```

op

Stark	2
Lannister	1
Greyjoy	1

* Read CASE from the doc

Like Operator

```
select * from emp where fname like '%on%'
```

→ '% on %' → any no. of chars before & after on.

- * LIKE '245%' → start with 245.
- * LIKE '246' → end with 246
- * LIKE '--n%' → 2 - skip first 2 chars before n.
- * LIKE 'a-n' → single char match a-n, adn, etc
- * LIKE '-A-T' → -SAT-, WATI, MATI

Escape statement in LIKE query

```

Select *
From tablename
where name like CONCAT('%.', @in-searchText, '%') ESCAPE '\'
    
```

→ prepend '\' to every char in search text so that results are correct

eg - in-searchText

a-b

a%.bcd-a

op without ESCAPE

X

X

with ESCAPE

✓

✓

Search for range of chars.

```
select * from emp where fname like '[A-F]%'
```

→ fname that starts with A to F.

Match by range or set

→ LIKE 'Fa-aJan'

→ any but not man

- LIKE '[lmnop]any' → many but not gany
- LIKE '[^a-g]any' → ^ → negation. many but not gany.
- LIKE '[^lmnop]any' → gany but not many

Wildcard Characters

SQL wildcards are —

- 1) % → zero or more chars.
- 2) _ → single char
- 3) [charlist] → set and range of chars to match [arg], [abc]
- 4) [^charlist] → Matches only a char not specified in brackets

filter results using WHERE and HAVING.

BETWEEN → is inclusive

Select * from sales where qty BETWEEN 10 and 17.
 $\underline{\text{qty} \geq 10 \text{ AND } \text{qty} \leq 17}$

→ date BETWEEN '2013-07-11' AND '2013-05-24'

→ Iname BETWEEN 'D' AND 'L' → Iname alphabetically falls b/w D & L.
 → 'A,B,C, M' → not included

Use HAVING with AGGREGATE functions.

Unlike WHERE clause, HAVING can be used with aggregate functions.

Where EXISTS.

Select * from tablename t WHERE EXISTS
 (Select 1 from tablename1 t1 where t1.id = t.id)

Select records in tablename that have records matching in tablename1.

OFFSET & LIMIT

Select * from tablename LIMIT 20, 20. -- offset, limit.
 (mysql)

EXCEPT

Select dataset except where values are in other dataset
 -- dataset schemas must be identical

SELECT 'Data1' as 'Column' UNION ALL
 SELECT 'Data2' as 'Column'

EXCEPT

SELECT 'Data2' as 'Column'

Returns data, ~~data~~.

Explain and Describe

Explain → shows you how the query will be executed
This shows if query uses an index or if you could optimize your query by adding an index.

Chapter 16: EXPLAIN and DESCRIBE

Section 16.1: EXPLAIN Select query

An Explain in front of a `select` query shows you how the query will be executed. This way you can see if the query uses an index or if you could optimize your query by adding an index.

Example query:

```
explain select * from user join data on user.test = data.fk_user;
```

Example result:

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	user	index	test	test	5	(null)	1	Using where; Using index
1	SIMPLE	data	ref	fk_user	fk_user	5	user.test	1	(null)

on `type` you see if an index was used. In the column `possible_keys` you see if the execution plan can choose from different indexes or if none exists. `key` tells you the actual used index. `key_len` shows you the size in bytes for one index item. The lower this value is the more index items fit into the same memory size and they can be faster processed. `rows` shows you the expected number of rows the query needs to scan, the lower the better.

Section 16.2: DESCRIBE tablename;

`DESCRIBE` and `EXPLAIN` are synonyms. `DESCRIBE` on a tablename returns the definition of the columns.

```
DESCRIBE tablename;
```

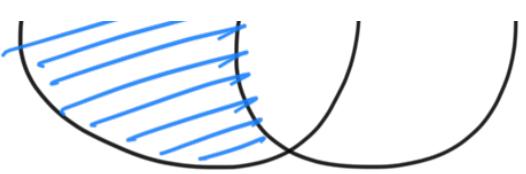
Example Result:

COLUMN_NAME	COLUMN_TYPE	IS_NULLABLE	COLUMN_KEY	COLUMN_DEFAULT	EXTRA
id	int(11)	NO	PRI	0	auto_increment
test	varchar(255)	YES		(null)	

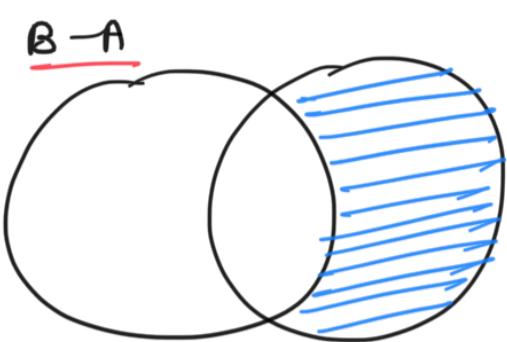
Here you see the column names, followed by the columns type. It shows if null is allowed in the column and if the column uses an index. The default value is also displayed and if the table contains any special behavior like an auto_increment.

MacBook Air



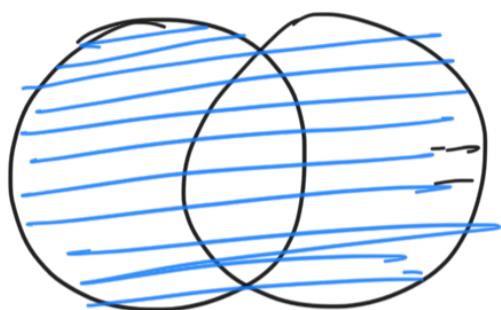


left JOIN
tableB B on A.key = B.key
where B.key IS NULL.

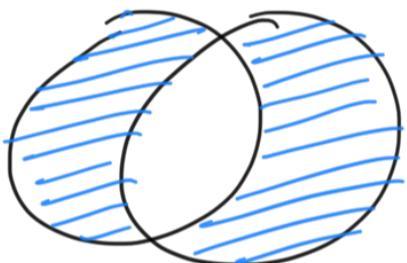


A.key IS NULL

FULL OUTER JOIN



Select <fields> from tableA A
FULL OUTER JOIN
tableB B on
A.key = B.key.



Select <fields> from tableA A
FULL OUTER JOIN
tableB B
ON A.key = B.key
where A.key IS NULL OR B.key IS NULL

Implicit Join → (not advised)

join using from, where clauses on tables
- does not contain JOIN command.

CROSS JOIN $(A \times B) \rightarrow$ Cartesian Product.
num in A x num in B \rightarrow table size

UPDATE

Update with data from other table

② Set Phonenum. for any employee who is also a customer
and currently does not have a phone number set.

UPDATE

Employees

SET phonenumber =

(select e.phonenumber FROM customers c WHERE
e.fname = c.fname
AND c.lname = e.lname)

WHERE Employee.phonenumber IS NULL.

Using INNER JOIN

```

UPDATE Employees
SET phonenumber = c.phonenumber
FROM Employees e
    INNER JOIN customer c
        ON e.fname = c.fname AND e.lname = c.lname
WHERE
    phonenumber is null

```

Modifying existing values.

```

Update Carr
Set totalcost = totalcost - 100
Where Id = 3 or Id = 4

```

Update specified rows

```

Update Carr
Set status = 'READY'
Where Id > 4

```

Update all rows

```

Update Carr
Set status = 'READY'

```

Capturing updated records.

Sometimes one wants to capture the records that have just been updated

```
CREATE TABLE #Tempupdated (ID INT)
```

```

Update table1 Set col1 = 42
    OUTPUT inserted.ID INTO #Tempupdated
    Where ID > 50 -

```

CREATE DATABASE

Create database database-name → creates empty db

Create Table

Create table from select

Create table clone AS select * from employees.

Create table clone AS

value of column name " " insert no following

```
select id, concat(firstname, ' ', lastname) as full_name  
from Employees  
where id > 10
```

Create New Table.

```
Create table tablename (  
    ID int identity(1,1) primary key not null,  
    fname varchar(10) not null,  
    lname varchar(10) not null,  
    number varchar(10),  
    DeptId int FOREIGN KEY REFERENCES Dept(Id)  
) ;
```

$\text{identity}(1,1) \Rightarrow$ col will have auto generated values starting at 1 and incrementing by 1 for each new row

Duplicate a table

```
CREATE TABLE newtable LIKE oldtable;  
INSERT INTO newtable Select * from oldtable;
```

creating a temporary or a in memory table

Create temp TABLE MyTable(...) → SQLite or Postgres.

SQL Server.

{
temp.
Create table #TempPhysical(...); → create a temp table local to the session
Create table ##TempPhysicalVisibleToEveryone(...)
→ visible to everyone.
DECLARE @TempMemory TABLE(...); → create an in memory table

CREATE FUNCTION

Argument	Description
function-name	name of fn.
list-of-parameters	params
return-data-type	return type
function-body	code
Scalar-expression	scalar value returned by function

```
CREATE FUNCTION firstword (@input varchar(1000))  
RETURNS varchar(1000)
```

AS

BEGIN

```
    DECLARE @Output varchar(1000)
```

```
    SET @Output = SUBSTRING (@input, 0, CHARINDEX
```

```

        WHEN 0 THEN LEN(@input)+1
        ELSE CHARINDEX(' ', @input)
    END
    RETURN @output
END

```

TRY/CATCH

Transaction in a TRY/CATCH

This will rollback both inserts due to invalid Datetime.

BEGIN TRANSACTION

BEGIN TRY

```

        INSERT INTO dbo.sde VALUES (8-2, GETDATE(), 1)
        _____
        COMMIT TRANSACTION

```

(---, not adate, 2)
invalid.

END TRY

BEGIN CATCH

```

        THROW
        ROLLBACK TRANSACTION.

```

END CATCH.

UNION / UNION ALL

- Used to combine two SELECT statement results without any duplicate.
- In order to use UNION both select statement should have same no. of columns with same datatype in same order. but the length of column can be different.

```

Select fname, lname
FROM HR-EMPLOYEES
WHERE pos = 'Mgr'
UNION ALL
SELECT fname, lname
FROM FIN-EMPLOYEE
WHERE pos = 'Mgr'

```

→ 2 ppl in diff dept with same
fname/lname
keep both.

UNION → removes duplicates

UNION ALL → does not remove duplicate

ALTER TABLE

- modify column/ constraint in a table

Add Column(s)

ALTER TABLE Employees

ADD StandardRate INT NULL

ITW STRUCTURE HAVE NOT NULL DEFAULT GETDATE(),
DOB DATE NULL

Drop column

```
ALTER TABLE emp  
DROP COLUMN columnname;
```

Add Primary Key

ALTER TABLE emp ADD pk-EmployeeID PRIMARY KEY (ID)
 |
 (ID, fname)

ALTER COLUMNS

ALTER TABLE EMPLOYEES

ALTER COLUMN StartDate DATETIME NOT NULL DEFAULT (GETDATE())

PROP constraint

ALTER TABLE Emp

DROP CONSTRAINT defaultsalary

Note : Ensure constraints of a column are dropped before dropping a column

INSERT

from another table

INSERT INTO Table1 Select * from table2;

INSERT NEW ROW

```
INSERT INTO CUSTOMERS VALUES (1, 'Zack', '221105')
```

Insert only specified columns

```
Insert into customer (name, contact)  
VALUES ('smith', '121034')
```

Insert multiple rows.

```
Insert intotbl-name(field1,field2,field3)  
VALUES (1,2,3),(4,5,6),(7,8,9);
```

MERGE

=> UPSERT \Rightarrow Update + Insert -

CROSS APPLY, OUTER APPLY

SELECT *

FROM Department D

CROSS APPLY C

SELECT *

→ Selects data from Department table & uses CROSS APPLY

```
FROM EMPLOYEE E  
WHERE E.DepartmentId = D.DeptId  
) A
```

GO

SELECT *

```
FROM DEPARTMENT D  
INNER JOIN Employee E  
ON E.DepartmentId = D.DeptId.
```

to evaluate Employee table for each record of Dept table.

→ join emp and ept and all matching records are produced.

SELECT *

```
FROM Department D  
OUTER APPLY (
```

SELECT *

```
FROM Employee E  
WHERE E.DepartmentId = D.DeptId.
```

) A

GO

SELECT *

```
FROM Department D
```

```
LEFT OUTER JOIN Employee E
```

```
ON D.DepartmentId = E.DepartmentId
```

GO

→ similar op
but execution plan will be different.

* when APPLY is required?

e.g.:

```
CREATE FUNCTION dbo.fnname (@DeptId AS int)
```

```
RETURNS TABLE
```

```
AS
```

```
RETURN (
```

```
SELECT * FROM Employee E  
where E.DeptId = @DeptId
```

```
)
```

GO

```
SELECT * FROM Department D
```

```
CROSS APPLY dbo.fnname(D.DepartmentId)
```

GO

```
SELECT * FROM Department D
```

```
OUTER APPLY dbo.fnname(D.DepartmentId)
```

GO

Can we use join here → NO

* If you replace CROSS OUTER APPLY with

↓ ↓
Innerjoin / left outerjoin

You will get error -

The multi-part identifier 'Employee' could not be bound.

This is because with JOIN the execution context of outer query is different from execution context of function (or a derived table) and you cannot bind a value/variable from outer query to the function as parameter.
Hence APPLY operator is used for such queries.

DELETE

Delete all rows (Truncate is better in performance as it ignores trigger and indexes and logs to just delete the data)

```
DELETE FROM Employees
```

Delete certain row with WHERE

```
DELETE FROM Emp WHERE ID > 10;
```

TRUNCATE → delete all rows and resets values such as autoincrement. It also doesn't log individual row deletion.

```
TRUNCATE TABLE Employees;
```

Delete

- 1) remove one or many (all) rows
- 2) row based operation
ie - each row is deleted
- 3) slower
- 4) does not reset autoninc.
- 5)

Truncate

- 1) removes all rows
- 2) Page Operation
Entire data page is reallocated.
- 3) faster way to delete all rows.
- 4) Reset AutoIncrement counters
- NO foreign keys must be present otherwise you will get an error.

DROP TABLE

check for existence before dropping

```
DROP TABLE IF EXISTS myTable;
```

Simple Drop

```
DROP TABLE tablename;
```

DROP/DELETE DATABASE

```
DROP DATABASE [dbo].[Employee]
```

-remove Employee db

* Always have a backup before dropping if required further

CASCADING DELETE

```

ALTER TABLE room
ADD CONSTRAINT PK-clientID FOREIGN KEY (RM-CL-ID)
REFERENCES CLIENT(ID)
ON DELETE CASCADE
ON UPDATE CASCADE

```

GRANT / REVOKE

```

REVOKE/GRANT SELECT, UPDATE
ON Employees
TO User1, User2

```

PRIMARY KEYS.

```

Create table Emp(
    Id INT NOT NULL,
    fname varchar(40),
    PRIMARY KEY (Id, fname)
);

```

Auto Increment

MySQL

```

CREATE TABLE Employees (
    Id int NOT NULL AUTO_INCREMENT,
    PRIMARY KEY (Id)
);

```

SQL Server

```

CREATE TABLE Employees(
    Id int NOT NULL IDENTITY,
    PRIMARY KEY (Id)
);

```

Indexes

- Data structure that contains pointers to the contents of a table arranged in a specific order, to help database optimize queries.
- eg: index of a book.
- When an index exists on a column used in query's WHERE, JOIN, ORDER BY, it can substantially improve query performance.

Create Index.

```
Create INDEX ix-car-emp ON cars (EmployeeId);
```

- creates index on employeeid in cars.

index can contain more than 1 column.

```
CREATE INDEX xindex ON cars(id,EmpId,Name)
```

- * The index will be useful in queries asking to sort or select by those columns

Sorted Index

If you use an index that is sorted the way you would retrieve it, the SELECT statement would not do additional sorting when in retrieval.

```
create index myscore ON Scoreboard (score DESC);
```

When you execute

```
select * from scoreboard order by score ASC
```

The db would not do additional sorting since the idx is in that order.

Partial or Filtered Index:

```
create index started-orders  
ON orders(productid)  
where order-state='started';
```

Drop, Disable, Rebuild an index

```
DROP INDEX index.name ON car;
```

* Recreating an index may be slow and computationally expensive.

* As an alternative the index can be disabled

```
ALTER INDEX cars-emp ON cars DISABLE;  
This allows the table to retain the structure,  
along with metadata about the index.
```

Critically, this retains index statistics, so it is possible to easily evaluate the change. If warranted the index can later be rebuilt instead of being recreated completely.

```
ALTER INDEX cars-emp ON cars REBUILD;
```

Critically, this retains the index...

Critically, this retains the index statistics, so that it is possible to easily evaluate the change. If warranted, the index can then later be rebuilt, instead of being recreated completely.

```
ALTER INDEX ix_cars_employee_id ON Cars REBUILD;
```

Section 37.5: Clustered, Unique, and Sorted Indexes

Indexes can have several characteristics that can be set either at creation, or by altering existing indexes.

```
CREATE CLUSTERED INDEX ix_clast_employee_id ON Employees(EmployeeId, Email);
```

The above SQL statement creates a new clustered index on Employees. Clustered indexes are indexes that dictate the actual structure of the table; the table itself is sorted to match the structure of the index. That means there can be at most one clustered index on a table. If a clustered index already exists on the table, the above statement will

GoodKicker.com - SQL Notes for Professionals 103

fail. (Tables with no clustered indexes are also called heaps.)

```
CREATE UNIQUE INDEX uq_customers_email ON Customers(Email);
```

This will create an unique index for the column Email in the table Customers. This index, along with speeding up queries like a normal index, will also force every email address in that column to be unique. If a row is inserted or updated with a non-unique Email value, the insertion or update will, by default, fail.

```
CREATE UNIQUE INDEX ix_eid_desc ON Customers(EmployeeID Desc);
```

This creates an index on Customers which also creates a table constraint that the EmployeeID must be unique. (This will fail if the column is not currently unique - in this case, if there are employees who share an ID.)

```
CREATE INDEX ix_eid_desc ON Customers(EmployeeID Desc);
```

This creates an index that is sorted in descending order. By default, indexes (in MSSQL server, at least) are ascending, but that can be changed.

Section 37.6: Rebuild Index

Over the course of time B-Tree indexes may become fragmented because of updating/deleting/inserting data. In SQLServer terminology we can have internal (index page which is half empty) and external (logical page order doesn't correspond physical order). Rebuilding index is very similar to dropping and re-creating it.

We can re-build an index with:

```
ALTER INDEX index_name REBUILD;
```

By default rebuilding index is offline operation which locks the table and prevents DML against it, but many RDBMS allow online rebuilding. Also, some DB vendors offer alternatives to index rebuilding such as REORGANIZE (SQL Server) or COALESCE/BHRINK SPACE(Oracle).

Section 37.7: Inserting with a Unique Index

```
UPDATE Customers SET Email = 'richard@123@example.com' WHERE id > 1;
```

This will fail if an unique index is set on the Email column of Customers. However, alternate behavior can be defined for this case:

```
UPDATE Customers SET Email = 'richard@123@example.com' WHERE id = 1 ON DUPLICATE KEY;
```

GROUP BY vs DISTINCT

Group by is used in a combination with aggregation functions.
DISTINCT is used to list a unique combination of
 distinct values for the specified columns.

e.g. storename user_id ordered

A	43	1
B	57	2
A	43	3
C	82	4
A	21	5

Select distinct storename, user_id from table;

Op

A	43
B	57
C	82
A	21

String Functions

- perform operations on string values and return either numeric or string values

Concatenate

Select 'Hello' || 'World' || '!' ; → returns HelloWorld!

not supported in SQL Server

SELECT CONCAT('H','W') → HW.

SELECT CONCAT('H','W','!') → HW! (Oracle does not support
CONCAT of more than
two strings)

Select 'Foo' + CAST(42 AS varchar(8)) + 'Bar';

- SQL Server < 2012

Length

- SELECT LEN('HELLO ') → 5

Does not count trailing spaces.

- SELECT DATALENGTH('HELLO ') → 6

counts trailing spaces.

Note - DataLength returns the length of underlying byte representation of string which depends on charset used to store string (ASCII, unicode etc)

e.g. DATALENGTH('Hello') = 5 (ASCII)
 = 10 (Unicode)

Oracle

Select length('Hello') from dual → 5

Trim empty spaces

```
SELECT LTRIM('HELO')  
--a - RTRIM('HELO')  
--b - LTRIM(RTRIM('HELO')) } MySQL
```

mysql & oracle
select TRIM (' hello ')
 LTrim
 RTrim

Upper and lower case

SELECT UPPER('Hello') → HELLO
← LOWER → hello

Split a string

Select value from STRING-SPLIT('Hi I am..', ' ');
→ split on space.

Replace

REPLACE(string to search, String to search for & replace,
String to replace with)

REPLACE ('Peter Steve Tom', 'Steve', 'Billy').
olp Peter Billy Tom.

REGEXP

Select 'bedded' REGEXP '[a-f]' → returns True
Select 'beam!' REGEXP '[a-f]' → returns False

substring .

SUBSTRINg(string-expression, start, length)

S&L strings are 1 indexed

SELECT SUBSTRING ('Hello', 1, 2) → 'He'
 y (y-3, 3) → 'lo'

LEFT, RIGHT

~~SELECT LEFT('Hello', 2) → He~~
Select RIGHT('Hello', 2) → lo

REVERSE

Select REVERSE ('Hello') → olleH

REPLICATE

Select REPLICATE ('Hello', 2) → HelloHello

Replace fn in SQL select and update query.

REPLACE (str, find, repl)

Select FirstName, REPLACE (Address, 'lSouth', 'lSouthern')
Address
From emp

We can also use it with update

INSTR

Return first occurrence of substring, return 0 if not found.

Select INSTR ('FOODarbar', 'Bar') → 4
→ 1 → ; 'Bar') → 0

PARENNAME

Returns specific part of given string (object name).

PARENNAME('obj.id.server.name', 2)
obj - server.

Functions (Aggregate)

conditional Aggregation -

Select customer,

sum(case when payment-type = 'credit' then amount else 0 end
as credit)

sum(case when pay-type = 'debit' then amount else 0 end)
as debit

from payments

group by customer

* List concatenation (later)

Aggregate fn's

sum(), min(), max(), avg(), count()

functions (Scalar/single row)

scalar fn → take one value as input and return one value as output for each row in result set.

Date / time -

<u>Datatype</u>	<u>format</u>
time	hh:mm:ss [.nnnnnnnn]
date	YYYY-MM-DD
smalldatetime	YYYY-MM-DD hh:mm:ss
datetime	YYYY-MM-DD hh:mm:ss [.nnnn]
datetime2	YYYY-MM-DD hh:mm:ss [.nnnnnnnnnn]
datetimeoffset	YYYY-MM-DD hh:mm:ss [.nnnnnnnn] [+/-] hh:mm

Datename() → returns name / value of specific part of date.

Select Datename(weekday, '2017-01-14') as Datename
Op Saturday

GetDate → determines current date and time of computer running the current SQL instance.
→ doesn't include the timezone difference.

Select GETDATE() as Systemdate.

Op - 2017-01-14 11:11:47.7230728

DATEDIFF → diff b/w two dates.

Select DATEDIFF (day, orderdate, shipdate)

You can have year, month, week, day, hr, min, s, ms.

DATEADD → add an interval to the date.

Select DATEADD (day, 20, '2017-01-14')

Op - 2017-02-03 00:00:00.000

Configuration and conversion function

@@SERVNAME - provides name of local server that is running SQL.

Select @@SERVNAME as 'server'

CAST / CONVERT → To perform any conversions that do not occur implicitly.

CAST(Hiredate as varchar(20)) As 'Cast' → 2003-02-04
use default style YYYY-MM-DD

CONVERT (varchar, HireDate, 3) As 'convert' → 09/02/03
- uses date & time style you specify. 3 → dd/mm/yy

Logical and Mathematical fn.

CHOOSE - returns an item from list of values based on position

Select CHOOSE (2, 'A', 'B', 'C') → Op - B.

IFF \rightarrow conditional-ternary $a = x ? a : b$
IFF (SaleryTD > 20000, 'A', 'B') as bonus.

\equiv SaleryTD > 20000 ? A : B.

SIGN

Select SIGN(-20) \rightarrow -1 -1 for negative
SIGN(20) \rightarrow +1 +1 for positive

POWER

Select POWER(3, 3) \rightarrow $3^{1^3} = 9$.
POWER(50, 3) \rightarrow $50^{1^3} = 125000$

Views - A view can filter some rows from base table or project only some columns from it.

```
Create view view-name AS  
Select id, fname, salary  
from Employees E  
where hiredate > date '2015-01-01';
```

```
Select * from view-name
```

Complex views

A view can be really complex query (aggregation, join, subquery etc) just be sure to add column names for everything you select.

```
Create view dept-income AS  
Select d.dname as Deptname, sum(e.salary)  
from Employees e  
JOIN Department d  
on e.deptid = d.id  
Group by d.dname;
```

Materialized views:

a view whose results are physically stored and must be periodically refreshed in order to remain current.

- Therefore useful in storing results of complex queries when realtime results are not required.
- Can be created in oracle & PostgreSQL.
- Other db offer similar functionality - SQL indexed views etc.

Section 48.1: PostgreSQL example

similar functionality, such as SQL Server's indexed views or DB2's materialized query tables.

Section 48.1: PostgreSQL example

```
CREATE TABLE mytable (number INT);
INSERT INTO mytable VALUES (1);

CREATE MATERIALIZED VIEW myview AS SELECT * FROM mytable;

SELECT * FROM myview;

number
-----
1
(1 row)

INSERT INTO mytable VALUES(2);

SELECT * FROM myview;

number
-----
1
2
(2 rows)

REFRESH MATERIALIZED VIEW myview;
    ^
SELECT * FROM myview;

number
-----
1
2
(2 rows)
```

Comments :

single line -

Select *

From emp -- This is a comment.

Multi line -

```

/* This is a
multiline comment */
Select *
from emp

```

foreign keys

- ensure data integrity by enforcing that values in one table must match values in other table.
- The datatype of fk must match datatype of referenced key.

A few tips using FK -

- * must reference a UNIQUE (or Primary) key in parent-table.
- * Entering a NULL value in FK does not raise an error
- * FK constraints can reference tables within same db
- * FK constraints can refer to another column in same table
(self reference)

Creating a table with fk

Create table Heropowers

```

ID INT NOTNULL ,
HeroId int references Superheroes(ID)
);

```

SEQUENCE

Create sequence orders_sq

Start with 1000

Increment by 1;

→ create a sequence with a starting value of 1000 and increment by 1.

General -

Create sequence seq-name

Start with initial-val

Increment by value

MINVALUE min → minimum value of sequence

MAXVALUE max → max → _____

CYCLE/ NO CYCLE; → cycle - when reach its set limit start from beginning.

nocycle → An exception will be thrown if exceeds max. value.

Using sequences

`seq-name.NEXTVAL` → returns next value in sequence.

can be used for Insert, Update, Select
`Insert into Orders values (orders-seq.NEXTVAL, 1032);`
`Select order-seq.NEXTVAL from dual;`

Update Orders

```
Set Order-ID = orders-seq.NEXTVAL  
where custid = 581;
```

Subqueries:

Subquery in FROM clause -

acts similar to a temporary table that is generated during the execution of the query and lost afterwards.

```
Select Managers.id, Employees.salary  
FROM C  
  Select id from Employees where ManagerId is NULL) As Managers  
JOIN Employees on Managers.id = Employees.id
```

Subquery in select clause.

```
Select id, fname, lname,  
  (Select count(*) from Cars where Cars.customerID =  
    customers.id) As NumCars  
FROM Customers
```

Subquery in where clause.

```
Select * from Employees  
where salary = (Select max(salary) from Employees)
```

Correlated Subqueries

Aka Synchronized or coordinated subqueries are nested queries that make references to the current row of outer query

```
Select EmployeeId  
FROM Employee As eOuter  
WHERE salary >  
  (Select Avg(salary)  
    FROM Employees eInner  
    ) where eInner.deptId = eOuter.deptId
```

Filter query results using query on different table

Select * from employees
where employeeId not in (Select empId from supervisor)
→ All employees not in supervisor table.

Execution blocks

Using BEGIN ... END

BEGIN

UPDATE Employee SET ph = '55723127' WHERE id = 1;
UPDATE emp SET salary = 1000 WHERE id = 2;

END

Stored Procedure

- A prepared SQL that you can save so that the code can be reused over and over again.

Create PROCEDURE procedure-name
@LastName nvarchar(50),
@FirstName nvarchar(50) } -parameters

AS

SELECT FirstName, LastName, Dept
FROM tablename
WHERE FirstName = @FirstName
AND LastName = @LastName
AND EndDate IS NULL.

Calling the procedure

EXECUTE procedure-name 'smith', 'jones';
or
EXEC procedure-name @LastName = 'Jones', @FirstName = 'Sm';

GO

TRIGGERS

- are stored programs which are automatically executed when some events occur.

Triggers can be on response of following event

DML - Delete, Insert, Update

DCL - Create, Alter, Drop

DB operation - Server error, Logon, Logoff, startup, shutdown

Benefits

- Generating some derived column values automatically

- Ensuring referential integrity
- event logging, and storing info on table access.
- Auditing
- Synchronous replication of tables
- Imposing security authorizations
- Prevent invalid transactions.

Syntax -

```

CREATE [OR REPLACE] TRIGGER trigger-name
{ BEFORE | AFTER | INSTEAD OF }
{ INSERT [OR] | UPDATE [OR] | DELETE }
[OF column]
ON table-name
[REFERENCING OLD AS o NEW AS n]
[FOR EACH ROW]
WHEN (condition)
DECLARE
    Declaration Statements
BEGIN
    Executable Statements
EXCEPTION
    Exception Handling Statements
END;

```

Eg -

```

CREATE OR REPLACE TRIGGER salary-chg
BEFORE DELETE OR INSERT OR UPDATE
ON CUSTOMERS
FOR EACH ROW
when (NEW.ID > 0)
DECLARE
    sal-diff-no;
BEGIN
    sal-diff-no := :NEW.SALARY - :OLD.SALARY ;
    dbms_output.put_line ('OldSal: ' || :OLD.SALARY);
END;

```

Eg -

```

create trigger mytrigger
    ON mytable
    AFTER INSERT
AS
BEGIN
    INSERT INTO myaudit (mytableid, user)
        (select mytableid, user-user from inserted)

```

END

Transactions

Simple Transaction

BEGIN TRANSACTION

```
INSERT INTO deletedemp (empid, date, user)
    (Select 123, getdate(), curr_user);
```

```
DELETE FROM emp WHERE id = 123;
```

COMMIT TRANSACTION.

Rollback Transaction

BEGIN TRY

BEGIN TRANSACTION

```
INSERT INTO users (id, name, age)
    VALUES (1, 'Bob', 24)
```

```
DELETE FROM users WHERE name = 'Todd'
```

COMMIT TRANSACTION

END TRY

BEGIN CATCH

ROLLBACK TRANSACTION

END CATCH

Savepoint - you can rollback to a savepoint

```
savepoint cravpoint_name;
```

```
rollback savepoint_name;
```

Autocommit

```
SET AUTOCOMMIT ON;
```

————— off;

SQL keywords are not case sensitive however it is common practice to write them in uppercase

SQL Injection

- is an attempt to access a website's database tables by injecting SQL in form field.
- If webserver does not protect against SQL injection attack, a hacker can trick the db into running SQL code which can access all user data, modify, delete db.

Userid : 105 OR 1=1

Select * from user where userid = 105 OR 1=1;

will give all users data.

Use SQL parameters for protection -

SQL engine checks params to ensure it is correct for its column and are treated literally and not as a part of SQL to be executed.

```
txtname = getRequestId("UserId");
txtSQL = "SELECT * FROM User WHERE UserId = @0";
db.execute (txtSQL, txtname)
```

PLSQL

PL/SQL is a combination of SQL along with procedural features of programming languages.

PL/SQL programming language was developed by

The PL/SQL programming language was developed by Oracle Corporation in the late 1980s as a procedural extension language for SQL and the Oracle relational database. Following are certain notable facts about PL/SQL –

- PL/SQL is a completely portable, high-performance transaction-processing language.
- PL/SQL provides a built-in, interpreted and OS independent programming environment.
- PL/SQL can also directly be called from the command-line **SQL*Plus interface**.
- Direct call can also be made from external programming language calls to database.
- PL/SQL's general syntax is based on that of ADA and Pascal programming language.
- Apart from Oracle, PL/SQL is available in **TimesTen in-memory database** and **IBM DB2**.

Features of PL/SQL

PL/SQL has the following features –

- PL/SQL is tightly integrated with SQL.
- It offers extensive error checking.
- It offers numerous data types.
- It offers a variety of programming structures.
- It supports structured programming through functions and procedures.
- It supports object-oriented programming.
- It supports the development of web applications and server pages.

Advantages of PL/SQL

PL/SQL has the following advantages –

- SQL is the standard database language and PL/SQL is strongly integrated with SQL. PL/SQL supports both static and dynamic SQL. Static SQL supports DML operations and transaction control from PL/SQL block. In Dynamic SQL, SQL allows embedding DDL

Advantages of PL/SQL

PL/SQL has the following advantages –

- SQL is the standard database language and PL/SQL is strongly integrated with SQL. PL/SQL supports both static and dynamic SQL. Static SQL supports DML operations and transaction control from PL/SQL block. In Dynamic SQL, SQL allows embedding DDL statements in PL/SQL blocks.
- PL/SQL allows sending an entire block of statements to the database which reduces network traffic and provides high performance for the application.
- PL/SQL gives high productivity to programmers as it can query, update and insert data in a database.
- PL/SQL saves time on design and debugging by strong features like exception handling, encapsulation, data hiding, and object-oriented data type declarations.
- Applications written in PL/SQL are fully portable.
- PL/SQL provides high security level.
- PL/SQL provides access to predefined SQL packages.
- PL/SQL provides support for Object-Oriented Programming.
- PL/SQL provides support for developing Web Applications and Server-side scripts.

PLSQL is a block structured language i.e - PLSQL programs are divided into logical blocks of code

Each block consists of three sub parts -

Declarations - DECLARE keyword.

- optional section, defines all variables, cursors, subprograms and other elements to be used in the program.

Executable commands -- BEGIN ... END. in b/w these.

- mandatory section.

- should have atleast 1 executable line of code which may be NULL to indicate nothing should be executed

Exception Handling → starts with EXCEPTION keyword.

- optional.

DECLARE

<declaration sections>

BEGIN.

<executable command(s)>

EXCEPTION

<exception handling>

END;

e.g. DECLARE

 message VARCHAR(20):='Hello World';

 BEGIN

 dbms_output.put_line(message);

 END;

/ → to run code from command line - you need to type / at beginning of first blank line after last line of code -

PLSQL Identifiers -

constants, variables, exceptions, procedures, cursors, reserved words.

- not case sensitive.

- can't use reserved keyword as identifier.

Delimiters : - a symbol with special meaning .

+ , -, *, / add, subtract, mult, divide

%, Attribute indicator

' char string delimiter

--- comment delimiters

	component selector.
(,)	Expression or list delimiter.
:	float variable indicator.
,	item separator
"	Quoted identifier delimiter.
=	Relational operator.
@	Remote alias indicator
;	Statement terminator
:=	Assignment operator
⇒	Association operator.
	concat
**	exponent
<<, >>	label delimiter (begin and end)
/*, */	multiline comment
--	singleline comment
..	range operator
<,>, <=,>=	relational operator
<=, !=, ~=, ^=	not equal to.

Data types

1) Scalar → single values with no internal components, such as number, Date, Boolean.

Numerics - PLS-Integer, Binary-Integer, Binary-Float, Binary Double
 Number(prec,scale), Decimal(prec,scale), Numeric(prec,scale)
 float, int, integer, smallint, real.

Character - char, varchar2, raw, nchar, nvarchar2, long, longraw, rowid, ~~rowwid~~.

Boolean - True, False, ~~Null~~

Datetime.

2) LOB - large object (text, image, video clips, sound waveforms).
 Bfile, Blob, Clob, NClob,

3) Composite : collections, records.

4) Reference : pointers to other data items

Null → nothing.

- can be assigned but cannot be equated with anything, including itself.

Variables:

variable-name [CONSTANT] datatype [NOT NULL] {:= | DEFAULT} initial-value
 eg pi CONSTANT double precision :=3.1415;
 eg name varchar(25)

operator precedence

$*$, $/$	expotent
$+, -, \sim$	identity, negation
$\ast, /$	mult, div
$+, -, \sim, //$	add, subtract, concat
	comparison

NOT

AND

OR

Loops

Basic Loop

Declare

x number := 10

BEGIN

LOOP

$x := x + 10$

IF $x > 50$ Then

exit;

END IF;

END LOOP;

END;

/

while loop

while condition loop

seq. of statements

END LOOP;

eg.

Declare a

a number := 10;

BEGIN

while $a < 20$ loop

$a := a + 1;$

END LOOP;

END;

For

FOR counter IN initial-val .. final-val LOOP
seq. of statements

END LOOP;

DECLARE

a number := 10

BEGIN

FOR a IN 10 .. 20 LOOP // FOR a in REVERSE 10..20 Loop
dbmns. output.println(a)

END LOOP;

END;

conditions

IF condition THEN

S;

END IF

eg IF ($a <= 20$) THEN

$c := c + 1$

END IF

IF CONDITION THEN

S;

—

```
    ELSE
        S2;
    END IF

if-then-elif
IF condition THEN
    S1;
ELSIF Condition2 THEN
    S2;
ELSIF cond3 THEN
    S3;
ELSE
    S4;
END IF;
```

Case

CASE selector

WHEN YOU THEN SI

WHEN val2 THEN s2;

- 1 -

ELSE **Sn**; -- default case.

END CASE;

String

Hello that's you → 'Hello that's you'

Arrays

Vairay

-each element has index, size can be changed dynamically.

Create or replace Type array-type-name IS Varray(n) of
valid attribute name.

Create or Replace TYPE namearray AS varray(3) of Varchar2)

8

Type namearray is varray(5) of integer;

DECLARE

Type namesarray is varray(5) of integer;
Mark namesarray;

BEGIHN

```

marks := array(1..12, 3, 4, 90, 95)
total := marks.count
for i in 1..total loop

```

FOR i in 1...total LOOP

```
END LOOP;  
END;  
/
```

functions and Procedures

return a single value, mainly used to compute and return a value	do not return a value directly. mainly used to perform an action
---	---

Procedure

```
create [or Replace] procedure procedure-name  
[ (parameter-name [IN|OUT|IN OUT] type [, ...]) ]  
} IS / AS {  
BEGIN  
    //Body  
END procedure-name;
```

AS - standalone procedures.

```
create or replace procedure greet
AS
BEGIN
    output putline('Hello')
END;
```

Call -

EXECUTE greet;
or

SEAN

greet ;
greet;

1

Drop procedure

~~DROP procedure procedure now;~~
~~||| greet;~~

function

```
create or replace function fn-name  
[(parameter-name [IN|OUT|INOUT] type [, ...])]  
RETURN      return-data-type  
{  
    IS | AS  
BEGIN
```

END [function-name];

Call -

BEGIN

C := totalcustomers()

END;

CURSORS

Oracle creates a memory area called Context area to process SQL statement, which contains all info needed for processing, eg. no. of rows processed etc.

A cursor is a pointer to context area.

PLSQL controls context area through cursor

-The set of rows the cursor holds is called active_set

There are 2 types of cursors

Implicit

- automatically created when SQL is executed.
- we can't control implicit cursor

SQL%rowcount = no of rows affected

Explicit

- programmer defined
- should be defined in declaration section of PLSQL Block.
- It is created on select statements which returns more than one row.

Explicit cursor

Declaring

```
CURSOR c-customers IS  
    select id, name, addr from customer;
```

Opening the cursor

```
OPEN c-customers.
```

Fetching the cursor, accessing one row at a time

```
FETCH c-customers INTO c-id, c-name, c-addr;
```

Closing the cursor

```
CLOSE c-customers;
```

DECLARE

```
DECLARE
    c_id customers.id%type;
    c_name customers.name%type;
    c_addr customers.address%type;
    CURSOR c_customers is
        SELECT id, name, address FROM customers;
BEGIN
    OPEN c_customers;
    LOOP
        FETCH c_customers into c_id, c_name, c_addr;
        EXIT WHEN c_customers%notfound;
        dbms_output.put_line(c_id || ' ' || c_name || ' ' || c_addr);
    END LOOP;
    CLOSE c_customers;
END;
/
```

Records - later

Exceptions - an error condition during a program execution.
System defined User defined.

Syntax

DECLARE
 ~~declaration section~~

BEGIN
 ~~executable cmd~~.

EXCEPTION
 WHEN exception1 THEN
 // handling statement
 WHEN exception2 THEN
 //

 :
 WHEN others THEN default
 // ...

Raising Exceptions

exceptions are raised automatically by db server but user can also raise.

DECLARE
 exception-name EXCEPTION; ↑ user defined exception
BEGIN
 IF condition THEN
 RAISE exception-name;
 END IF
EXCEPTION
 WHEN exception-name THEN
 Statement;
END;

Datetime

ADD_MONTHS(x, y) Add y months to x

LAST_DAY(x) → last day of month

MONTHS_BETWEEN(x, y)

NEXT_DAY(x, day) - returns datetime of next day after x.
New time.

`getdate()` → current datetime
`getdate() -> el →`
`CURRENT_TIMESTAMP() ->`

Normalization → reduce redundancy

- eliminate insert, update, delete anomalies.
- structure logically.

FULL NAMES

FULL NAMES	PHYSICAL ADDRESS	MOVIES
Janet Jones	First Street Plot No 4	Pirates of the Caribbean the Tita
Robert Phil	3 rd Street 34	Forgetting Marshall Little Gi
Robert Phil	5 th Avenue	Clash of

1NF - Atomic (no multivalued attributes)

- each cell should contain single value
- each record needs to be unique

FULL NAMES

FULL NAMES	PHYSICAL ADDRESS	MOVIES RENTED
Janet Jones	First Street Plot No 4	Pirates of the Caribbean
Janet Jones	First Street Plot No 4	Clash of the Tita
Robert Phil	3 rd Street 34	Forgetting Sarah Marshall
Robert Phil	3 rd Street 34	Daddy's Little Gi
Robert Phil	5 th Avenue	Clash of the Tita

$2NF \rightarrow 1NF$ $(Stid, coun, countee) \rightarrow dependency$
 → No functional dependency.
 → Single column PK that does not functionally depend on any subset of candidate key relation

Rule 1- Be in 1NF

2NF (Second Normal Form) Rules

- Rule 1- Be in 1NF
- Rule 2- Single Column Primary Key that does not functionally dependant on any subset of candidate key relation

It is clear that we can't move forward to make our simple database in 2nd Normalization form unless we partition the table above.

MEMBERSHIP ID	FULL NAMES	PHYSICAL ADDRESS	SALUTATION
1	Janet Jones	First Street Plot No 4	Ms.
2	Robert Phil	3 rd Street 34	Mr.
3	Robert Phil	5 th Avenue	Mr.

MEMBERSHIP ID	MOVIES RENTED
1	Pirates of the Caribbean
1	Clash of the Titans
2	Forgetting Sarah Marshal
2	Daddy's Little Girls
3	Clash of the Titans

We have divided our 1NF table into two tables viz. Table 1 and Table2. Table 1 contains member information. Table 2 contains information on movies rented.

We have introduced a new column called Membership_id which is the primary key for table 1. Records can be uniquely identified in Table 1 using membership id

Database – Foreign Key

Transitive dependency- changing a non key column may cause any of the other non key column to change

Id	Name	Salutation
1	Robert	Mr
2	Janet	Ms
3	Robert	Mr

Rebeca —————> Mrs.

$A \rightarrow B$ } transitive dependency.
 $B \rightarrow C$ }
 name \rightarrow state
 state \rightarrow country.
 $st, state \rightarrow State, country$.

3NF \rightarrow 2NF

→ no transitive dependency.

NF Example

NF Example

Below is a 3NF example in SQL database:

MEMBERSHIP ID	FULL NAMES	PHYSICAL ADDRESS	SALUTATION ID
1	Janet Jones	First Street Plot No 4	2
2	Robert Phil	3rd Street 34	1
3	Robert Phil	5th Avenue	1

MEMBERSHIP ID	MOVIES RENTED
1	Pirates of the Caribbean
1	Clash of the Titans
2	Forgetting Sarah Marshall
2	Daddy's Little Girls
3	Clash of the Titans

SALUTATION ID	SALUTATION
1	Mr.
2	Ms.
3	Mrs.
4	Dr.

We have again divided our tables and created a new table which stores Salutations.

There are no transitive functional dependencies, and hence our table is in 3NF

In Table 3 Salutation ID is primary key, and in Table 1 Salutation ID is foreign to primary key in Table 3

Now our little example is at a level that cannot further be decomposed to attain higher normal form types of normalization in DBMS. In fact, it is already in higher normalization forms. Separate efforts for moving into ne-

β CNF - Boyce Codd Normal Form. (2.5NF).

Even when DB is in 3NF there would still be anomalies if it has more than one candidate key

Keys -

- 1 Superkey - combo of keys.
- 2 Candidate key - set of cols that uniquely identify a record.
- 3 PK - no nulls, unique
- 4 FK
- 5 Alternate key - candidate key, not selected as PK currently
- 6 Composite key - eg - PK(id, name)
- 7 Unique key - 1 null allowed, unique

Possible Candidate Keys

ID	RollNo	Name	EnrollNo	Address	DeptID	
1	60391401	Avin	AX101	Delhi	1	FK_Relationship
2	60391402	Amit	AX102	Noida	1	Possible Candidate
3	60391403	Mohan	AX103	Ghaziabad	2	
4	60391404	Pavan	AX104	Ghaziabad	2	
5	60391405	Deepak	AX105	Delhi	3	
6	60391406	Jitendra	AX106	Delhi	3	Foreign Key

copyright dotnet-tricks.com

Primary Key **Alternate Keys** **Unique Key**

Super Key